

Libraries ¶

```
In [2]: 1 # Suppressing all warnings
2 import warnings
3 warnings.filterwarnings("ignore")
4
5 # Multiprocessing
6 import multiprocessing
7
8 # Threading
9 import threading
10
11 # Time
12 import time
13
14 # Data handling
15 import pandas as pd
16 import numpy as np
17
18 # Plotting
19 import matplotlib.pyplot as plt
20 import seaborn as sns
21
22 # Oversampling
23 from imblearn.over_sampling import ADASYN
24
25 # Metrics
26 from sklearn import metrics
27
28 # Preprocessing
29 from imblearn.over_sampling import ADASYN
30 from sklearn.model_selection import train_test_split
31 from sklearn.decomposition import PCA
32
33 # Models
34 from sklearn.pipeline import Pipeline
35 from sklearn.svm import SVC
36 from sklearn.neighbors import KNeighborsClassifier
37 from sklearn.tree import DecisionTreeClassifier
38 from sklearn.linear_model import LogisticRegression
39 from xgboost import XGBClassifier
40 from sklearn.ensemble import AdaBoostClassifier
41 # from catboost import CatBoostClassifier
42
43 # Evaluation metrics
44
45 # Cross validation
46 from sklearn.model_selection import RandomizedSearchCV, GridSearchCV, train_test_split
47 from skopt import BayesSearchCV
48 from sklearn.metrics import classification_report
```

```
In [2]: 1 def multi_time():
2         print("Starting thread...")
3         start = time.time()
4
5         while True:
6             elapsed = int(time.time() - start)
7             print(f"Time elapsed: {elapsed} seconds", end = "\r")
8             time.sleep(1)
```

Loading CSV file onto Pandas DataFrame

```
In [74]: 1 df = pd.read_csv("player_stats.csv", na_values="None")
2         df = df[["X", "Y", "xG", "h_a", "situation", "shotType", "lastAction", "result"]]
3         # df = df.Loc[:50000, :].copy()
```

Check for missing data and percentage

```
In [75]: 1 print("Percentage of data missing from each column\n")
2         for col in df.columns:
3             if len(df[df[col].isnull() == True]):
4                 print(f"{col}: {len(df[df[col].isnull() == True])/len(df)*100}%",
5                       end = "")
6                 print(f" or {len(df[df[col].isnull() == True])} rows")
7             else:
8                 print(f"{col}: {0}%")
```

Percentage of data missing from each column

X: 0%
Y: 0%
xG: 0%
h_a: 0%
situation: 0%
shotType: 0%
lastAction: 9.228053426520521% or 10246 rows
result: 0%

Dropping rows with missing values

```
In [76]: 1 df = df[df["lastAction"].isnull() == False].reset_index(drop = True).copy()
2         df = df.drop(df[df["result"]=="OwnGoal"].index)
3         df = df[df.situation != 'Penalty']
```

Encoding categorical columns

- h_a: One-Hot Encoder
- situation: One-Hot Encoder
- shotType: One-Hot Encoder
- lastAction: One-Hot Encoder
- result: One-Hot Encoder

In [77]:

```
1 print(f"\x1b[32mUnique values in:\x1b[0m {df['h_a'].value_counts()}")
2 print(f"\x1b[32mUnique values in:\x1b[0m {df['situation'].value_counts()}")
3 print(f"\x1b[32mUnique values in:\x1b[0m {df['shotType'].value_counts()}")
4 print(f"\x1b[32mUnique values in:\x1b[0m {df['lastAction'].value_counts()}")
5 print(f"\x1b[32mUnique values in:\x1b[0m {df['result'].value_counts()}")
```

```
Unique values in: h      54108
a      44615
Name: h_a, dtype: int64
Unique values in: OpenPlay      82980
FromCorner      8514
SetPiece      3918
DirectFreekick      3311
Name: situation, dtype: int64
Unique values in: RightFoot      48933
LeftFoot      32717
Head      16725
OtherBodyPart      348
Name: shotType, dtype: int64
Unique values in: Pass      41618
Cross      15324
Aerial      7000
Chipped      6630
TakeOn      6616
Rebound      4705
Throughball      3515
Standard      3311
BallRecovery      3081
HeadPass      2644
BallTouch      1919
LayOff      707
Dispossessed      547
Tackle      221
CornerAwarded      194
BlockedPass      171
Foul      159
Goal      88
Interception      65
End      64
OffsidePass      23
GoodSkill      23
Challenge      18
Card      18
SubstitutionOn      16
Clearance      13
OffsideProvoked      9
KeeperPickup      7
Save      7
FormationChange      5
Start      3
ChanceMissed      1
ShieldBallOpp      1
Name: lastAction, dtype: int64
Unique values in: MissedShots      34019
SavedShot      25463
BlockedShot      24329
Goal      12795
ShotOnPost      2117
Name: result, dtype: int64
```

```

In [78]: 1 remapping = {"Goal": 1,
2               "MissedShots": 0,
3               "SavedShot": 0,
4               "BlockedShot": 0,
5               "ShotOnPost": 0,
6               }
7
8 lastAction = {"Cross": ["Aerial", "Cross", "Chipped"] ,
9              "Pass":["Pass", "Throughball", "HeadPass", "BallTouch", "TakeOn"],
10             "Dispossessed":["Dispossessed", "Tackle", "BlockedPass", "Interception",
11                             "Challenge", "Clearance", "BallRecovery"],
12             "Other":["Standard", "Rebound", "LayOff", "CornerAwarded",
13                     "Foul", "Goal", "End", "OffsidePass",
14                     "GoodSkill", "Card", "SubstitutionOn", "OffsideProvoked",
15                     "Save", "KeeperPickup",
16                     "FormationChange", "Start", "ChanceMissed" , "ShieldBallOpp"]}
17
18 df["result"] = df["result"].map(remapping)
19 for k, v in lastAction.items():
20     for i in v:
21         df["lastAction"].replace(to_replace = i, value = k, inplace = True)
22
23 df_encoded = pd.get_dummies(df,
24                             drop_first = True)
25 df_encoded = df_encoded.loc[:50000, :]

```

```

In [79]: 1 X = df_encoded.drop(["result"], axis = 1).copy()
2        Y = df_encoded["result"].copy()

```

```

In [80]: 1 X_train, X_test, y_train, y_test_1 = train_test_split(X, Y,
2                                                         train_size = 0.8,
3                                                         stratify = Y)

```

Step 2: Grid Search

```

In [81]: 1 pipe = Pipeline([("classifier", LogisticRegression())])
2
3 param_grid = [
4     {
5         'classifier': [LogisticRegression(random_state = 42)],
6         'classifier__solver': ['saga', 'lbfgs'],
7         'classifier__C': [0.1, 1],
8     },
9     {
10        'classifier': [SVC(random_state = 42)],
11        'classifier__C': [0.1, 1],
12        'classifier__degree': [1, 2],
13        'classifier__gamma': [0.1, 1],
14    },
15    {
16        'classifier': [XGBClassifier(random_state = 42)],
17        'classifier__n_estimators': [10, 1000],
18        'classifier__max_depth': [10, 50],
19    },
20    {
21        'classifier': [AdaBoostClassifier()],
22        'classifier__base_estimator': [DecisionTreeClassifier(random_state = 42)],
23        'classifier__learning_rate': [0.01, 0.1, 1],
24        'classifier__base_estimator__criterion': ['gini', 'entropy'],
25        'classifier__base_estimator__max_depth': [1, 5, 9],
26    }
27 ]
28
29 bs = GridSearchCV(pipe, param_grid, cv = 5, verbose = 3)
30 bs.fit(X_train, y_train)

[CV 5/5] END classifier=LogisticRegression(random_state=42), classifier__C=1, classifier__solver=saga; score=0.876 total time= 0.1s
[CV 1/5] END classifier=LogisticRegression(random_state=42), classifier__C=1, classifier__solver=lbfgs; score=0.879 total time= 0.1s
[CV 2/5] END classifier=LogisticRegression(random_state=42), classifier__C=1, classifier__solver=lbfgs; score=0.879 total time= 0.1s
[CV 3/5] END classifier=LogisticRegression(random_state=42), classifier__C=1, classifier__solver=lbfgs; score=0.875 total time= 0.1s
[CV 4/5] END classifier=LogisticRegression(random_state=42), classifier__C=1, classifier__solver=lbfgs; score=0.877 total time= 0.1s
[CV 5/5] END classifier=LogisticRegression(random_state=42), classifier__C=1, classifier__solver=lbfgs; score=0.876 total time= 0.1s
[CV 1/5] END classifier=SVC(random_state=42), classifier__C=0.1, classifier__degree=1, classifier__gamma=0.1; score=0.867 total time= 37.1s
[CV 2/5] END classifier=SVC(random_state=42), classifier__C=0.1, classifier__degree=1, classifier__gamma=0.1; score=0.867 total time= 37.5s
[CV 3/5] END classifier=SVC(random_state=42), classifier__C=0.1, classifier__degree=1, classifier__gamma=0.1; score=0.867 total time= 38.7s
[CV 4/5] END classifier=SVC(random_state=42), classifier__C=0.1, classifier__degree=1, classifier__gamma=0.1; score=0.867 total time= 37.6s

```

```

In [82]: 1 bs.best_params_

```

```

Out[82]: {'classifier': AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=1,
                                         random_state=42),
                                         learning_rate=1),
          'classifier__base_estimator': DecisionTreeClassifier(max_depth=1, random_state=42),
          'classifier__base_estimator__criterion': 'gini',
          'classifier__base_estimator__max_depth': 1,
          'classifier__learning_rate': 1}

```

```
In [83]: 1 y_preds_1 = bs.predict(X_test)
          2 acc = metrics.accuracy_score(y_test_1, y_preds_1)
```

```
In [84]: 1 cr_1 = classification_report(y_test_1, y_preds_1)
          2 print(classification_report(y_test_1, y_preds_1))
```

	precision	recall	f1-score	support
0	0.89	0.98	0.93	8502
1	0.60	0.20	0.30	1307
accuracy			0.88	9809
macro avg	0.74	0.59	0.62	9809
weighted avg	0.85	0.88	0.85	9809

Recall isn't good for minority class

This can be attributed to the heavy class imbalance in the dataset. We can try to balance this by first filtering data and then using ADASYN

Higher cutoff for number of shots per player

```
In [85]: 1 def preprocessing(df, num_rows = 50000):
2
3     print("\x1b[32mStarting preprocessing...\x1b[31m")
4
5     df = df[["X", "Y", "xG", "h_a", "situation", "shotType", "lastAction", "result"]]
6     df = df[df["lastAction"].isnull() == False].reset_index(drop = True).copy()
7     df = df.drop(df[df["result"] == "OwnGoal"].index)
8     df = df[df.situation != 'Penalty']
9     remapping = {"Goal": 1,
10                  "MissedShots": 0,
11                  "SavedShot": 0,
12                  "BlockedShot": 0,
13                  "ShotOnPost": 0,
14                  }
15
16     lastAction = {"Cross": ["Aerial", "Cross", "Chipped"],
17                  "Pass": ["Pass", "Throughball", "HeadPass", "BallTouch", "TakeOn"],
18                  "Dispossessed": ["Dispossessed", "Tackle", "BlockedPass",
19                                   "Interception", "Challenge", "Clearance",
20                                   "BallRecovery"],
21                  "Other": ["Standard", "Rebound", "LayOff", "CornerAwarded",
22                             "Foul", "Goal", "End", "OffsidePass",
23                             "GoodSkill", "Card", "SubstitutionOn", "OffsideProvoked",
24                             "Save", "KeeperPickup", "FormationChange",
25                             "Start", "ChanceMissed", "ShieldBallOpp"]}
26
27     print("Remapping target column")
28     df["result"] = df["result"].map(remapping)
29
30     print("Remapping lastAction column")
31     for k, v in lastAction.items():
32         for i in v:
33             df["lastAction"].replace(to_replace = i, value = k, inplace = True)
34
35     print("One hot encoding remaining columns")
36     df_encoded = pd.get_dummies(df,
37                                 drop_first = True)
38
39     if num_rows != None:
40         print(f"Using only {num_rows} rows")
41         df_encoded = df_encoded.loc[:num_rows, :]
42     else:
43         print(f"Using all rows")
44
45     print("Preprocessing done\x1b[0m")
46     return df_encoded
```


In [86]:

```
1 def search(X_train, y_train, verbosity = 1, use_svc = True):
2
3     print("\x1b[32m Starting search...\x1b[31m")
4
5     print("Building base pipeline")
6     pipe = Pipeline([("classifier", LogisticRegression())])
7
8     print("Building parameter grid")
9     if use_svc:
10         print("Using SVC")
11         param_grid = [
12             {
13                 'classifier': [LogisticRegression(random_state = 42)],
14                 'classifier__solver': ['saga', 'lbfgs'],
15                 'classifier__C': [0.1, 1],
16             },
17             {
18                 'classifier': [SVC(random_state = 42)],
19                 'classifier__C': [0.1, 1],
20                 'classifier__degree': [1, 2],
21                 'classifier__gamma': [0.1, 1],
22             },
23             {
24                 'classifier': [XGBClassifier(random_state = 42)],
25                 'classifier__n_estimators': [10, 1000],
26                 'classifier__max_depth': [10, 50],
27             },
28             {
29                 'classifier': [AdaBoostClassifier()],
30                 'classifier__base_estimator': [DecisionTreeClassifier(random_state = 42)],
31                 'classifier__learning_rate': [0.01, 0.1, 1],
32                 'classifier__base_estimator__criterion': ['gini', 'entropy'],
33                 'classifier__base_estimator__max_depth': [1, 5, 9],
34             }
35         ]
36     else:
37         print("Replacing SVC with KNN")
38         param_grid = [
39             {
40                 'classifier': [LogisticRegression(random_state = 42)],
41                 'classifier__solver': ['saga', 'lbfgs'],
42                 'classifier__C': [0.1, 1],
43             },
44             {
45                 'classifier': [KNeighborsClassifier()],
46                 'classifier__n_neighbors': [3, 4, 5],
47                 'classifier__p': [1, 2],
48             },
49             {
50                 'classifier': [XGBClassifier(random_state = 42)],
51                 'classifier__n_estimators': [10, 1000],
52                 'classifier__max_depth': [10, 50],
53             },
54             {
55                 'classifier': [AdaBoostClassifier()],
56                 'classifier__base_estimator': [DecisionTreeClassifier(random_state = 42)],
57                 'classifier__learning_rate': [0.01, 0.1, 1],
58                 'classifier__base_estimator__criterion': ['gini', 'entropy'],
59                 'classifier__base_estimator__max_depth': [1, 5, 9],
60             }
61         ]
62     print("Creating Grid Search")
63     bs = GridSearchCV(pipe, param_grid, cv = 5, verbose = verbosity)
```

```
64     print("Fitting train data onto Grid Search")
65     bs.fit(X_train, y_train)
66
67     print(f"\x1b[32mBest Parameters:\x1b[0m\n {bs.best_params_}")
68
69     return bs
```

In [87]:

```
1  def test(best_model, X_test, y_test):
2      y_pred = best_model.predict(X_test)
3      acc = metrics.accuracy_score(y_test, y_pred)
4
5      print(f"\x1b[32mClassification report on test accuracy:\x1b[0m\n {classification_r
6
7      return y_pred, classification_report(y_test, y_pred)
```

In [88]:

```
1 df2 = pd.read_csv("player_stats.csv", na_values="None")
2
3 df3 = df2.copy()
4
5 print("Before filtering, df shape", df3.shape)
6
7 df3_grouped = df3.groupby(["player", "result"]).size()
8 df3['shot_count'] = df3.groupby('player')['player'].transform('size')
9 df4 = df3.groupby(["player", 'shot_count', "result"])[['shot_count']].count()
10 df4 = df4.rename(columns = {'shot_count': 'count'})
11 df5 = df4.sort_values(by = df4.index.names[1], ascending = False)
12 df5["result"] = df5.index.get_level_values(2)
13 df5.index = df5.index.droplevel(2)
14 df5['shot_count'] = df5.index.get_level_values(1)
15 df5.index = df5.index.droplevel(1)
16 df5 = df5[df5['shot_count'] > 40]
17 df3 = df3[df3["player"].isin(df5.index)]
18 print("After filtering, df shape", df3.shape)
19
20 df_encoded = preprocessing(df3, num_rows = None)
21 X = df_encoded.drop(["result"], axis = 1).copy()
22 Y = df_encoded["result"].copy()
23
24 X_train, X_test, y_train, y_test_2 = train_test_split(X, Y,
25                                                         train_size = 0.8,
26                                                         stratify = Y)
27
28 best_model_2 = search(X_train, y_train, verbosity = 3,
29                       use_svc = True)
30
31 y_preds_2, cr_2 = test(best_model_2, X_test, y_test_2)
32
33
```

1.6s

[CV 3/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=0.1; score=0.880 total time= 1

1.7s

[CV 4/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=0.1; score=0.878 total time= 1

1.8s

[CV 5/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=0.1; score=0.881 total time= 1

2.0s

[CV 1/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=1; score=0.878 total time= 11.9

s

[CV 2/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=1; score=0.879 total time= 11.6

Metrics didn't increase as much.

Let's try oversampling using ADASYN and replace KNN with SVC

In [29]:

```
1 df2 = pd.read_csv("player_stats.csv", na_values="None")
2
3 df3 = df2.copy()
4
5 print("Before filtering, df shape", df3.shape)
6
7 df3_grouped = df3.groupby(["player", "result"]).size()
8 df3['shot_count'] = df3.groupby('player')['player'].transform('size')
9 df4 = df3.groupby(["player", 'shot_count', "result"])[['shot_count']].count()
10 df4 = df4.rename(columns = {'shot_count': 'count'})
11 df5 = df4.sort_values(by = df4.index.names[1], ascending = False)
12 df5["result"] = df5.index.get_level_values(2)
13 df5.index = df5.index.droplevel(2)
14 df5['shot_count'] = df5.index.get_level_values(1)
15 df5.index = df5.index.droplevel(1)
16 df5 = df5[df5['shot_count'] > 40]
17 df3 = df3[df3["player"].isin(df5.index)]
18 print("After filtering, df shape", df3.shape)
19
20 df_encoded = preprocessing(df3, num_rows = 50000)
21 df5_g = df_encoded[df_encoded["result"] == 1]
22 df5_ng = df_encoded[df_encoded["result"] == 0]
23
24 print("\x1b[32mBefore ADASYN\x1b[0m")
25 print(f"Ratio of goals to not goals: {len(df5_g)/len(df5_ng)}")
26
27 print("\x1b[31mOversampling using ADASYN\x1b[0m")
28 X = df_encoded.drop(["result"], axis = 1).copy()
29 Y = df_encoded["result"].copy()
30
31 X_resampled, y_resampled = ADASYN().fit_resample(X, Y)
32 print("\x1b[32mAfter ADASYN\x1b[0m")
33 print(f"Number of samples added by ADASYN {len(y_resampled) - len(Y)}, or {(len(y_resampled) - len(Y)) / len(Y)}")
34 print(f"Ratio of goals to not goals: {len(y_resampled[y_resampled == 1]) / len(y_resampled)}")
35
36 X_train, X_test, y_train, y_test_3 = train_test_split(X_resampled, y_resampled,
37                                                         train_size = 0.8,
38                                                         stratify = y_resampled)
39
40 best_model_3 = search(X_train, y_train, verbosity = 3, use_svc = False)
41
42 y_preds_3, cr_3 = test(best_model_3, X_test, y_test_3)
```

```
ase_estimator__max_depth=5, classifier__learning_rate=0.1; score=0.787 total time= 13.6s
[CV 3/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=0.1; score=0.792 total time= 13.7s
[CV 4/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=0.1; score=0.831 total time= 14.3s

[CV 5/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=0.1; score=0.768 total time= 14.5s
[CV 1/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=1; score=0.817 total time= 14.5s
[CV 2/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeClassifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__base_estimator__max_depth=5, classifier__learning_rate=1; score=0.817 total time= 14.5s
```

Entire dataset since it can run faster

In [93]:

```
1 df2 = pd.read_csv("player_stats.csv", na_values="None")
2
3 df3 = df2.copy()
4
5 print("Before filtering, df shape", df3.shape)
6
7 df3_grouped = df3.groupby(["player", "result"]).size()
8 df3['shot_count'] = df3.groupby('player')['player'].transform('size')
9 df4 = df3.groupby(["player", 'shot_count', "result"])[['shot_count']].count()
10 df4 = df4.rename(columns = {'shot_count': 'count'})
11 df5 = df4.sort_values(by = df4.index.names[1], ascending = False)
12 df5["result"] = df5.index.get_level_values(2)
13 df5.index = df5.index.droplevel(2)
14 df5['shot_count'] = df5.index.get_level_values(1)
15 df5.index = df5.index.droplevel(1)
16 df5 = df5[df5['shot_count'] > 40]
17 df3 = df3[df3["player"].isin(df5.index)]
18 print("After filtering, df shape", df3.shape)
19
20 df_encoded = preprocessing(df3, num_rows = None)
21 df5_g = df_encoded[df_encoded["result"] == 1]
22 df5_ng = df_encoded[df_encoded["result"] == 0]
23
24 print("\x1b[32mBefore ADASYN\x1b[0m")
25 print(f"Ratio of goals to not goals: {len(df5_g)/len(df5_ng)}")
26
27 print("\x1b[31mOversampling using ADASYN\x1b[0m")
28 X = df_encoded.drop(["result"], axis = 1).copy()
29 Y = df_encoded["result"].copy()
30
31 X_resampled, y_resampled = ADASYN().fit_resample(X, Y)
32 print("\x1b[32mAfter ADASYN\x1b[0m")
33 print(f"Number of samples added by ADASYN {len(y_resampled) - len(Y)}, or {(len(y_resampled) - len(Y)) / len(Y)}")
34 print(f"Ratio of goals to not goals: {len(y_resampled[y_resampled == 1]) / len(y_resampled)}")
35
36 X_train, X_test, y_train, y_test_4 = train_test_split(X_resampled, y_resampled,
37                                                         train_size = 0.8,
38                                                         stratify = y_resampled)
39
40 best_model_4 = search(X_train, y_train, verbosity = 3, use_svc = False)
41
42 y_preds_4, cr_4 = test(best_model_4, X_test, y_test_4)
```

```
assifier(random_state=42), classifier__base_estimator__criterion=gini, classifier__base
__estimator__max_depth=9, classifier__learning_rate=1; score=0.894 total time= 39.2s
[CV 1/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeCl
assifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__b
ase_estimator__max_depth=1, classifier__learning_rate=0.01; score=0.674 total time=
7.5s
[CV 2/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeCl
assifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__b
ase_estimator__max_depth=1, classifier__learning_rate=0.01; score=0.680 total time=
7.7s
[CV 3/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeCl
assifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__b
ase_estimator__max_depth=1, classifier__learning_rate=0.01; score=0.677 total time=
7.7s
[CV 4/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeCl
assifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__b
ase_estimator__max_depth=1, classifier__learning_rate=0.01; score=0.681 total time=
8.0s
[CV 5/5] END classifier=AdaBoostClassifier(), classifier__base_estimator=DecisionTreeCl
assifier(random_state=42), classifier__base_estimator__criterion=entropy, classifier__b
```

Classification Reports

In [96]:

```
1 print("\x1b[32mBaseline evaluation metrics:\x1b[0m")
2 print("\x1b[31mNumber of rows = 50000")
3 print("List of models:\nLogistic Regression\nSVC\nXgboost\nAdaBoost\x1b[0m")
4 print(cr_1)
5
6 print("\x1b[36mMinority class recall and precision was low.\x1b[0m")
7 print("\x1b[36mChoosing top 75% percentile of players with respect to number of shots\x1b[0m")
8
9 print("\x1b[32mEvaluation metrics after choosing cutoff of number of shots for each player\x1b[0m")
10 print("\x1b[31mNumber of rows = 50000")
11 print("List of models:\nLogistic Regression\nSVC\nXgboost\nAdaBoost\x1b[0m")
12 print(cr_2)
13
14 print("\x1b[36mMetrics did not change by much\x1b[0m")
15
16 print("\x1b[36mOversampling using ADASYN to balance classes and replacing SVC with KNN\x1b[0m")
17 print("\x1b[32mEvaluation metrics after oversampling and replacing SVC with KNN:\x1b[0m")
18 print("\x1b[31mNumber of rows = 50000")
19 print("List of models:\nLogistic Regression\nKNN\nXgboost\nAdaBoost\x1b[0m")
20 print(cr_3)
21
22 print("\x1b[36mMetrics increased by a significant amount with respect to minority class\x1b[0m")
23 print("\x1b[36mRunning gridsearch again with entire dataset\x1b[0m")
24 print("\x1b[31mNumber of rows = entire dataset")
25 print("List of models:\nLogistic Regression\nKNN\nXgboost\nAdaBoost\x1b[0m")
26 print(cr_4)
```


Baseline evaluation metrics:

Number of rows = 50000

List of models:

Logistic Regression

SVC

Xgboost

AdaBoost

	precision	recall	f1-score	support
0	0.89	0.98	0.93	8502
1	0.60	0.20	0.30	1307
accuracy			0.88	9809
macro avg	0.74	0.59	0.62	9809
weighted avg	0.85	0.88	0.85	9809

Minority class recall and precision was low.

Choosing top 75% percentile of players with respect to number of shots

Evaluation metrics after choosing cutoff of number of shots for each player(top 75% percentile):

Number of rows = 50000

List of models:

Logistic Regression

SVC

Xgboost

AdaBoost

	precision	recall	f1-score	support
0	0.89	0.98	0.93	16711
1	0.60	0.22	0.32	2498
accuracy			0.88	19209
macro avg	0.75	0.60	0.63	19209
weighted avg	0.86	0.88	0.85	19209

Metrics did not change by much

Oversampling using ADASYN to balance classes and replacing SVC with KNN

Evaluation metrics after oversampling and replacing SVC with KNN:

Number of rows = 50000

List of models:

Logistic Regression

KNN

Xgboost

AdaBoost

	precision	recall	f1-score	support
0	0.88	0.93	0.90	8498
1	0.92	0.87	0.90	8494
accuracy			0.90	16992
macro avg	0.90	0.90	0.90	16992
weighted avg	0.90	0.90	0.90	16992

Metrics increased by a significant amount with respect to minority class and grid search ran computationally much faster.

Running gridsearch again with entire dataset

Number of rows = entire dataset

List of models:

Logistic Regression

KNN

Xgboost

AdaBoost

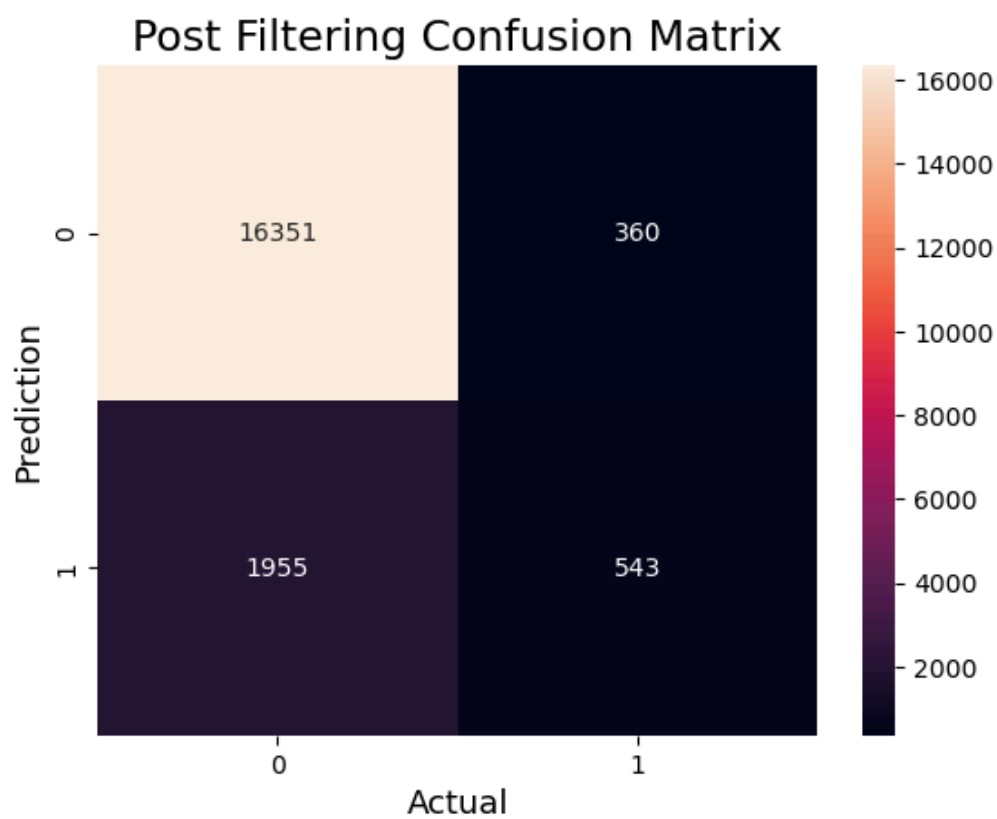

	precision	recall	f1-score	support
--	-----------	--------	----------	---------


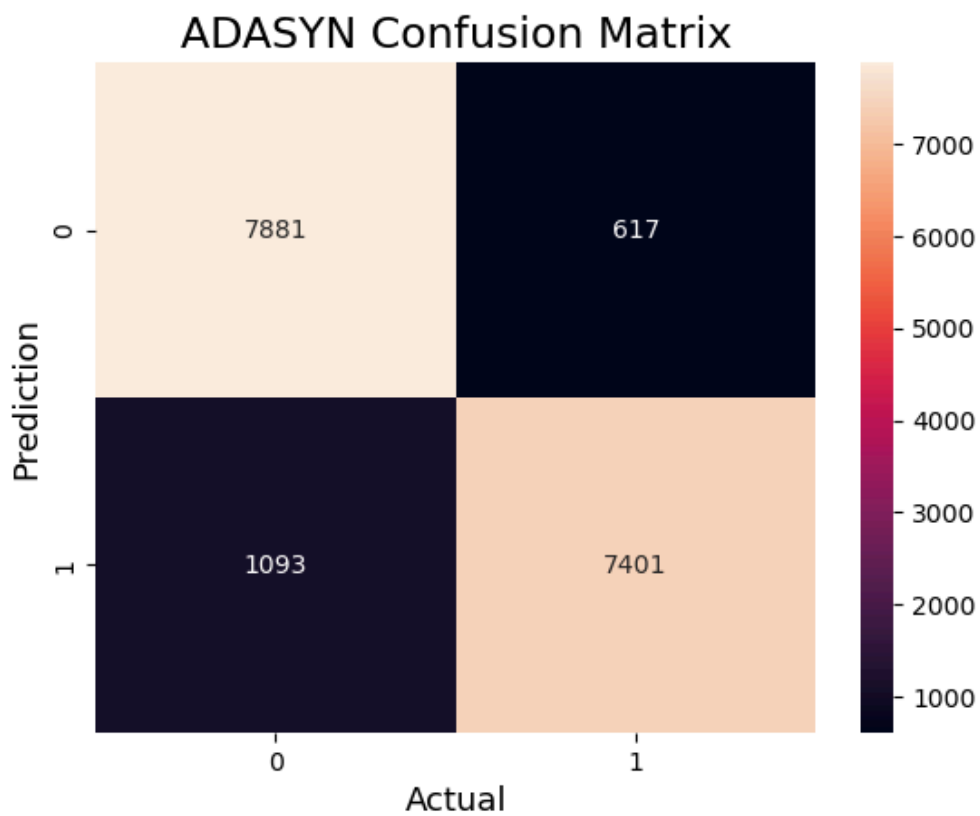
	0	0.89	0.93	0.91	16711
	1	0.93	0.89	0.91	16977
accuracy				0.91	33688
macro avg		0.91	0.91	0.91	33688
weighted avg		0.91	0.91	0.91	33688

Confusion Matrices

In [97]:

```
1 cm_1 = metrics.confusion_matrix(y_test_1, y_preds_1)
2 cm_2 = metrics.confusion_matrix(y_test_2, y_preds_2)
3 cm_3 = metrics.confusion_matrix(y_test_3, y_preds_3)
4 cm_4 = metrics.confusion_matrix(y_test_4, y_preds_4)
5
6 sns.heatmap(cm_1,
7             annot=True,
8             fmt='g',
9             xticklabels=range(2),
10            yticklabels=range(2))
11 plt.ylabel('Prediction',fontsize=13)
12 plt.xlabel('Actual',fontsize=13)
13 plt.title('Baseline Confusion Matrix',fontsize=17)
14 plt.show()
15
16 sns.heatmap(cm_2,
17             annot=True,
18             fmt='g',
19             xticklabels=range(2),
20            yticklabels=range(2))
21 plt.ylabel('Prediction',fontsize=13)
22 plt.xlabel('Actual',fontsize=13)
23 plt.title('Post Filtering Confusion Matrix',fontsize=17)
24 plt.show()
25
26 sns.heatmap(cm_3,
27             annot=True,
28             fmt='g',
29             xticklabels=range(2),
30            yticklabels=range(2))
31 plt.ylabel('Prediction',fontsize=13)
32 plt.xlabel('Actual',fontsize=13)
33 plt.title('ADASYN Confusion Matrix',fontsize=17)
34 plt.show()
35
36 sns.heatmap(cm_4,
37             annot=True,
38             fmt='g',
39             xticklabels=range(2),
40            yticklabels=range(2))
41 plt.ylabel('Prediction',fontsize=13)
42 plt.xlabel('Actual',fontsize=13)
43 plt.title('Final Confusion Matrix',fontsize=17)
44 plt.show()
```






AUC_ROC Curves

In [98]:

```
1 fpr_1, tpr_1, _ = metrics.roc_curve(y_test_1, y_preds_1)
2 auc_1 = metrics.roc_auc_score(y_test_1, y_preds_1)
3
4 plt.plot(fpr_1, tpr_1, label='Baseline ROC\nAUC = '+ str(auc_1))
5 # plt.plot([0, 1], [0, 1], '--', color='grey', label='Worst case')
6 plt.plot()
7 plt.xlabel('False Positive Rate')
8 plt.ylabel('True Positive Rate')
9 plt.title('ROC Curve')
10 # plt.legend();
11
12 fpr_2, tpr_2, _ = metrics.roc_curve(y_test_2, y_preds_2)
13 auc_2 = metrics.roc_auc_score(y_test_2, y_preds_2)
14
15 plt.plot(fpr_2, tpr_2, label='Filtered ROC\nAUC = '+ str(auc_2))
16 # plt.plot([0, 1], [0, 1], '--', color='grey', label='Worst case')
17 plt.plot()
18 plt.xlabel('False Positive Rate')
19 plt.ylabel('True Positive Rate')
20 plt.title('ROC Curve')
21 # plt.legend();
22
23 fpr_3, tpr_3, _ = metrics.roc_curve(y_test_3, y_preds_3)
24 auc_3 = metrics.roc_auc_score(y_test_3, y_preds_3)
25
26 plt.plot(fpr_3, tpr_3, label='ADASYN ROC\nAUC = '+ str(auc_3))
27 # plt.plot([0, 1], [0, 1], '--', color='grey', label='Worst case')
28 plt.plot()
29 plt.xlabel('False Positive Rate')
30 plt.ylabel('True Positive Rate')
31 plt.title('ROC Curve')
32 # plt.legend();
33
34 fpr_4, tpr_4, _ = metrics.roc_curve(y_test_4, y_preds_4)
35 auc_4 = metrics.roc_auc_score(y_test_4, y_preds_4)
36
37 plt.plot(fpr_4, tpr_4, label='Final ROC\nAUC = '+ str(auc_4))
38 plt.plot([0, 1], [0, 1], '--', color='grey', label='Worst case')
39 plt.plot()
40 plt.xlabel('False Positive Rate')
41 plt.ylabel('True Positive Rate')
42 plt.title('ROC Curve')
43 plt.legend(bbox_to_anchor = (1.02, 1.02));
```




Heatmap

Predicted Label Heatmap

In [94]:

```
1  # from highlight_text import fig_text
2  import matplotlib.path_effects as path_effects
3
4  from matplotlib.colors import LinearSegmentedColormap
5  from scipy.ndimage import gaussian_filter
6
7  from mplsoccer import Pitch, VerticalPitch, FontManager, Sboopen
8  from sklearn.model_selection import train_test_split
9
10 mean_with_zero = lambda x: np.sum(x) if not np.sum(np.sum(x)) else 0
11 pred_heatmap = pd.concat([pd.DataFrame(X_test), pd.Series(y_preds_4, index=X_test.index)
12
13 pred_heatmap = pred_heatmap.rename(columns = {0 : 'Result'})
14
15 plt.figure(dpi=4800)
16 path_eff = [path_effects.Stroke(linewidth=0.5, foreground='black'),
17             path_effects.Normal()]
18 pitch = VerticalPitch(pad_bottom = 1, half = True, goal_type = 'line', goal_alpha = 0.5,
19                      pitch_type = 'custom', pitch_length = 99.5, pitch_width = 100,
20                      line_zorder=2, line_color='white', corner_arcs=True, pitch_color=
21
22 fig, ax = pitch.draw(figsize=(15,15))
23 fig.set_facecolor('#22312b')
24
25 bin_statistic = pitch.bin_statistic(pred_heatmap.X * 100, pred_heatmap.Y * 100,
26                                     values = pred_heatmap.Result, statistic='sum', bins=
27
28 bin_statistic['statistic'] = gaussian_filter(bin_statistic['statistic'], -10)
29
30 pcm = pitch.heatmap(bin_statistic, ax=ax, cmap='plasma')
31 cbar = fig.colorbar(pcm, ax=ax, shrink=0.6)
32
33 cbar.outline.set_edgecolor('#efefef')
34 cbar.ax.yaxis.set_tick_params(color='#efefef')
35
36 ticks = plt.setp(plt.getp(cbar.ax.axes, 'yticklabels'), color='#efefef')
37 # labels = pitch.label_heatmap(bin_statistic, color='black', fontsize=10,
38 #                               ax=ax, ha='center', va='center',
39 #                               str_format='{:.2f}%', path_effects=path_eff)
```

<Figure size 30720x23040 with 0 Axes>




True Label Heatmap

In [95]:

```
1  # from highlight_text import fig_text
2  import matplotlib.path_effects as path_effects
3
4  from matplotlib.colors import LinearSegmentedColormap
5  from scipy.ndimage import gaussian_filter
6
7  from mplsoccer import Pitch, VerticalPitch, FontManager, Sboopen
8  from sklearn.model_selection import train_test_split
9
10 mean_with_zero = lambda x: np.sum(x) if not np.sum(np.sum(x)) else 0
11 pred_heatmap = pd.concat([pd.DataFrame(X_test), pd.Series(y_test_4, index=X_test.index)
12
13 # pred_heatmap = pred_heatmap.rename(columns = {0 : 'Result'})
14
15 plt.figure(dpi=4800)
16 path_eff = [path_effects.Stroke(linewidth=0.5, foreground='black'),
17             path_effects.Normal()]
18 pitch = VerticalPitch(pad_bottom = 1, half = True, goal_type = 'line', goal_alpha = 0.5,
19                      pitch_type = 'custom', pitch_length = 99.5, pitch_width = 100,
20                      line_zorder=2, line_color='white', corner_arcs=True, pitch_color=
21
22 fig, ax = pitch.draw(figsize=(15,15))
23 fig.set_facecolor('#22312b')
24
25 bin_statistic = pitch.bin_statistic(pred_heatmap.X * 100, pred_heatmap.Y * 100,
26                                     values = pred_heatmap.result, statistic='sum', bins=
27
28 bin_statistic['statistic'] = gaussian_filter(bin_statistic['statistic'], -10)
29
30 pcm = pitch.heatmap(bin_statistic, ax=ax, cmap='plasma')
31 cbar = fig.colorbar(pcm, ax=ax, shrink=0.6)
32
33 cbar.outline.set_edgecolor('#efefef')
34 cbar.ax.yaxis.set_tick_params(color='#efefef')
35
36 ticks = plt.setp(plt.getp(cbar.ax.axes, 'yticklabels'), color='#efefef')
37 # labels = pitch.label_heatmap(bin_statistic, color='black', fontsize=10,
38 #                               ax=ax, ha='center', va='center',
39 #                               str_format='{:.2f}%', path_effects=path_eff)
```

<Figure size 30720x23040 with 0 Axes>



In []:

1